

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/341710282>

Introduction to MongoDB

Presentation · February 2015

DOI: 10.13140/RG.2.2.13674.95687

CITATIONS

0

READS

198

1 author:



Mayur Patil
MIT Academy of Engineering
14 PUBLICATIONS 3 CITATIONS

SEE PROFILE

Some of the authors of this publication are also working on these related projects:

Project

Intrusion Detection by Forensic Method in Private Cloud using Eucalyptus [View project](#)

Introduction to



Prof Mayur S Patil,
Assistant Professor,
Dept. of Computer Engineering,
MIT Academy of Engineering, Alandi (D.)

Contents

- **Databases**
- **Purpose of Databases**
- **Types of Databases**
- **Relational Model**
- **Relational Databases**
- **What is MongoDB?**
- **Basics of MongoDB**
- **CRUD Terms**
- **Limitations**
- **Current Trending**

Databases

- Data + Base
- Add
- Access
- Update
- Delete

Purpose of Databases

- Easy to inject and retrieve data
- Able to store and use sets of data
- To get faster and accurate access of data through organized ways.
- Query data in a database (ask it questions).
- Relate data from two different tables together using JOINS.
- Create meaningful reports from data in a database.
- Information of a given type is always stored only once.
- Fault-tolerant.
- Concurrent; multiple users can use them at the same time without corrupting the data.

Applications:

Early Phase: Airline, Accounting, Railway reservation.

Now Trending:

- Social Networks: Facebook, Twitter.
- Search Engines: Bing, Google.
- Cloud Services: Amazon, Openstack.
- Big Data: Hadoop.

Types of Databases

- Relational Databases:

- MySQL – Facebook, Twitter.
- PostgreSQL -
- SQLite – Browsers.

- Non Relational Databases:

- MongoDB document
- CouchDB document
- ArangoDB
- Cassandra

Relational Model

- The **relational model (RM)** for database management is an approach to managing data using a structure and language.
- In the relational model of a database, all data is represented in terms of tuples and grouped into relations.
- A database organized in terms of the relational model is a relational database.

Relational Databases

- A database structured to recognize relations between stored items of information.
- Data organization approach
 - Tables
 - Row
 - Columns

RDBMS

- Introduction
 - Systems that used relational databases known as RDBMS.
 - It used SQL.
 - **DML – Update.**
 - **DDL – New things**
 - **DCL – Control Access**
- Advantages
 - Building a low-volume, medium-complexity suite of applications that will evolve over time.
 - Data volumes (duplicated) would be ridiculous if you didn't do a reasonable amount of normalization.
 - You simply don't see a cost/benefit advantage to moving away from proven legacy technology.
- Limitations
 - Data Complexity
 - Broken Keys and Records
 - Developer Expertise
 - Hardware Performance

Introduction to MongoDB

- MongoDB (from "humongous" i.e. huge + Monstrous) is a scalable, high-performance, open source, schema-free, document-oriented database.
 - mongodb.org
- It used NoSQL mechanism.
- A record in MongoDB is a document.

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```



Diagram illustrating the structure of a MongoDB document (a JSON-like object). The document is shown as a collection of field-value pairs. Arrows point from the text "field: value" to the corresponding field and value in the document.

Field	Value
name	"sue"
age	26
status	"A"
groups	["news", "sports"]

- The advantages of using documents are:
 - Documents (i.e. objects) correspond to native data types in many programming languages.
 - Data structure composed of field and value pairs.

Comparing Terminologies

SQL Terminologies	NoSQL Terminologies
Database	Database
Table	Collection
Row	Document

Features of MongoDB

- High Performance
 - I/O results
 - Faster keys
- High Availability
 - Replica Set
- Automatic Scaling
 - Sharding
 - Automatic balancing for changes in load and data distribution
 - Easy addition of new machines without down time
 - Scaling to one thousand nodes
 - No single points of failure
 - Automatic failover

- One or more shards, **each shard holds a portion of the total**
- **Each shard is backed by a replica set**
- Failure Management
- One or more **routers**, each one acts as a server for one or more clients.
- One or more clients, each one is (part of) the user's application and issues commands to a router via the mongo client library (driver) for its language.
- ***mongod*** is server process and mongos is router process.

Basics of MongoDB

- **JSON**

- Stores data in Object format.
- It supports nested looping i.e. objects within array.
- Documents enclosed in circular braces.
- Objects are enclosed in curly braces and separated by commas
- JSON supports
 - Number,
 - Strings,
 - Objects
 - Arrays

Example

```
db.things.save
```

```
(
```

```
{
```

```
  a : 1, b : 1,
```

```
    fruit:
```

```
    ["apple", "grapes", "pear" ]
```


```
}
```

```
)
```


BSON

- Its binary JSON i.e. format in which mongodb stores its data.
 - MDB drivers send and receive data as BSON from app.
 - On app side, MDB drivers maps BSON to native Data types in relative to programming language.
-
- Why BSON:
 - Lightweight: Its space required to for data keeping is minimum
 - Traversable: Writing and reading indexes to MDB
 - Efficient: Encoding data to BSON and Decoding from BSON by drivers for app.

Functioning of MongoDB



Working of MongoDB

- MongoDB is a **server process** that runs on Linux, Windows and OS X
 - It can be run both as a 32 or 64-bit application.
- Clients **connect** to the MongoDB process.
- MongoDB stores its data in files (default location is /data/db/), and uses **memory mapped files** for data management for efficiency.

CRUD Terms

SQL operation	MongoDB operations
Create	Create
Insert	Read
Update	Update
Delete	Remove

Limitations

- **Joins:** between two collections because it is difficult to scale out.
- **Transaction:** because documents are stored in hierarchical manner so it is not possible to access those items atomically.
- **Naming Restrictions**
 - Database Name Case Sensitivity
- **BSON Documents**
 - BSON Document Size
- **Max connection number is hardcoded to 20k.**

- Auto rollback is not for more than 300 MB; more than this manual intervention is needed.
- To shard a collection, it must be smaller than 256 GB, or else it will likely fail to shard.
- Map / Reduce are single-threaded.
- Map / Reduce cannot output to sharded collections.

Trending

- **Adobe – AEM**
- **Nokia – For Entertainment databases**
- **Facebook – Customer Data**
- **Google – Google Cloud Platform**
- **Ebay – Media Metadata**

THANKS !